BayCauRETM: R package for Bayesian Causal Inference for Recurrent Event Outcomes

11 September 2025

Summary

Observational studies often estimate the effect of medical treatments on the rate of a recurrent event within a specific follow-up window. Recurrent events (e.g., repeated hospitalizations, relapses) may occur multiple times during follow-up. Causal analysis is challenging because: 1) Recurrent events are jointly observed with a terminal event (e.g., death), which truncates future recurrences. 2) Both event counts and the terminal process are unobserved under dropout/censoring. 3) Patients initiate treatment at different times, yielding as many strategies as initiation times. Finally, 4) Treatments are not randomized, so formal causal methods are required to adjust for observed confounders.

This paper presents `BayCauRETM`, an `R` package for estimating causal effects of different treatment initiation strategies on a recurrent event outcome in the presence of death and censoring. Users specify an initiation time and supply a `data.frame` containing confounders and columns for censoring, death, and interval-specific recurrent counts, then define terminal and recurrent models via standard `R` formula syntax. Given these inputs, `BayCauRETM` performs causal adjustment and outputs adjusted expected event rates over follow-up under the specified initiation time. The package also provides diagnostic and visualization utilities.

Intended users include statisticians, epidemiologists, and health-services researchers analyzing observational data.

Statement of need

Standard software for time-to-event and recurrent-event data remains useful descriptively but generally does not target causal estimands while addressing complexities (1)-(4) above (Ghosh and Lin 2002; Schaubel and Zhang 2010; Janvin et al. 2024). Oganisian et al. (2024) developed Bayesian statistical methods that accommodate these complexities and conducted a thorough simulation-based validation of these methods. However, due to the focus on methodological development and validation, only proof-of-concept replication code was provided along with the paper. There is need for user-friendly, off-the-shelf software with readable help files that can implement the methods developed in Oganisian et al. (2024). `BayCauRETM` fills this methodological and practical gap by operationalizing the Bayesian approach in `R`. Its syntax is familiar to base-R users and mirrors standard regression functions such as `lm()` and `glm`, with extensive help pages accessible via `help()`. Thus, `BayCauRETM` provides the first user-friendly software for analyzing complex recurrent-event data while handling complexities (1)-(4) described in the Summary section above.

Data structure, model, and outputs

In this section, we provide an overview of the expected input data structure, models that are run under-the-hood, and expected outputs. We refer readers to Oganisian et al. (2024) for methodological details.

Data structure and preprocessing

The package expects longitudinal data in long, person-interval format. For follow-up time τ, the window $[0, \tau)$ is partitioned into K equal-length intervals $I_k = [\tau_{k-1}, \tau_k)$ for $k = 1, \dots, K$ with $\tau_0 = 0$ and $\tau_K = \tau$. Each row represents a patient-interval; a subject has one row per interval at risk.

Required `data.frame` variables are: subject ID, interval index k, treatment indicator (0 until the initiation interval, then 1), interval-specific count of recurrent events, and a terminal-event indicator (0 up to death, 1 thereafter). Optional variables include baseline covariates and lagged history (e.g., one-interval lag of the event count).

Each row contains a monotone death indicator at the start of interval k, T_k, a monotone treatment indicator by the end of interval k, A_k, the interval count Y_k and baseline covariates $L \in \mathcal{L}$.

Causal estimand and potential outcomes

Let $a(s) = (\underbrace{0, \dots, 0}_{s-1}, 1, \dots, 1)$ be the strategy that initiates treatment at interval $s \in \{1, 2, \dots, K+1\}$. Let $T_k^{a(s)}$ and $Y_k^{a(s)}$ denote, respectively, the death indicator and number of recurrent events that would have been observed in interval k under strategy $a(s)$. Formally, $T_k^{a(s)}$ is the potential death indicator following sequence $a(s)$ up to interval k, and $Y_k^{a(s)}$ is the corresponding potential number of events.

The target is the difference in average potential incidence rates over follow-up under two initiation times:

$$\Delta(s, s') = \mathbb{E}\left[\frac{\sum_{k=1}^{K} Y_k^{a(s)}}{K - \sum_{k=1}^{K} T_k^{a(s)}}\right] - \mathbb{E}\left[\frac{\sum_{k=1}^{K} Y_k^{a(s')}}{K - \sum_{k=1}^{K} T_k^{a(s')}}\right].$$

Model specification

The package runs a pair of discrete-time models conditional on shared treatment and covariate terms.

Here and throughout, we use overbar notation to denote the full history of the recurrent event process up to the previous interval, i.e., $\bar{Y}_{k-1} = (Y_1, Y_2, \dots, Y_{k-1})$.

1. Discrete-time hazard model for the terminal event that models death at a given interval conditional on survival up to that interval:
$$\lambda_k(a_k, \bar{Y}_{k-1}, l) = \Pr(T_k = 1 \mid T_{k-1} = 0, a_k, \bar{Y}_{k-1}, l).$$

2. Distribution for the number of event occurrences in a given interval conditional on survival through that interval:
$$f(y_k \mid a_k, \bar{Y}_{k-1}, l) = \Pr(Y_k = y_k \mid T_k = 0, a_k, \bar{Y}_{k-1}, l).$$

Here, $f(y_k \mid a_k, \bar{Y}_{k-1}, l)$ denotes the Poisson probability mass function with conditional mean (intensity) $\mu_k(a_k, \bar{Y}_{k-1}, l) = \mathbb{E}[Y_k \mid A_k, \bar{Y}_{k-1}, L]$. Together, these two models multiply to form a joint model for the terminal and recurrent event occurrence at a given interval.

The functions in `BayCauRETM` implement the following models for the hazard and intensity, respectively:

$$\text{logit}\,\lambda_k(a_k, \bar{Y}_{k-1}, l) = \beta_{0k} + l^\top \beta_L$$
$$+ y_{k-1}\beta_Y + \beta_A a_k,$$
$$\log \mu_k(a_k, \bar{Y}_{k-1}, l) = \theta_{0k} + l^\top \theta_L + y_{k-1}\theta_Y + \theta_A a_k.$$

The time-varying intercepts $\{\beta_{0k}\}$ and $\{\theta_{0k}\}$ parameterize the baseline hazard and event intensity, respectively. They are assigned a first-order autoregressive (AR1) smoothing prior. See Oganisian et al. (2024) for more details.

Posterior inference and g-computation

BayCauRETM conducts full posterior inference for the joint models using Stan (Carpenter et al. 2017) through the `rstan` interface, since the posterior distribution is not available in closed form. Stan is a probabilistic programming language that implements Hamiltonian Monte Carlo to generate posterior draws.

For each parameter draw obtained from Stan, `BayCauRETM` simulates the joint death-recurrent process under $a(s)$ and $a(s')$ to obtain a posterior draw of $\Delta(s, s')$. Reporting over many draws yields posterior samples of $\Delta(s, s')$, as described by Oganisian et al. (2024). The posterior mean and the 95% credible interval (2.5th and 97.5th percentiles) are reported.

Quickstart

Below we provide a minimal, copy-pastable example illustrating the core functionality of `BayCauRETM`: fitting a joint recurrent-event and terminal-event model and estimating causal effects under alternative treatment initiation strategies. A small example dataset is shipped with the package, allowing users to run the example without any external data dependencies.

Installation and setup

We first install the development version of the package from GitHub and load the required libraries. For reproducibility, we also set a random seed.

```r
# install.packages("pak")
pak::pak("LnnnnYW/BayCauRETM")

library(BayCauRETM)
library(dplyr)
library(tidyr)

set.seed(123)
```

Data loading

The package ships with a small example dataset used in the demonstration code. The dataset is stored in the package directory and can be loaded using `system.file()`. Loading this file creates a data frame named `df` in the workspace.

```r
# Load the example dataset shipped with the package
rdata_path <- system.file("demo_code", "data.Rdata", package = "BayCauRETM")
stopifnot(file.exists(rdata_path))
load(rdata_path)  # loads an object named `df`
```

Minimal preprocessing

The data are assumed to be in long, person-interval format. For a fast illustrative run, we subset the data to a small number of subjects and perform minimal preprocessing. This includes ordering observations by subject and time, constructing a discrete time index, creating a lagged event-count variable, removing incomplete cases, and standardizing continuous covariates. These steps mirror the preprocessing required for real applications but are kept intentionally simple here.

```r
# Minimal preprocessing
df_fit <- df %>%
  filter(id %in% 1:50) %>%                    # subset for quickstart
  arrange(id, k) %>%
  mutate(k_fac = as.integer(factor(k, levels = sort(unique(k))))) %>%
  group_by(id) %>%
  mutate(lagYk = if ("lagYk" %in% names(.)) replace_na(lagYk, 0) else lag(Yk, default = 0)) %>%
  ungroup() %>%
  drop_na(Tk, Yk, Ak, L.1, L.2) %>%
  mutate(
    L.1 = as.numeric(scale(L.1)),
    L.2 = as.numeric(scale(L.2))
  )

K <- length(unique(df_fit$k_fac))
```

Model fitting

We next fit the joint Bayesian model for the recurrent-event and terminal-event processes using the main
function `fit_causal_recur()`. The user specifies the outcome models through standard R formula syntax,
along with the relevant column names for subject ID, time index, treatment, and lagged history. Here we
use a single core and suppress verbose output for speed.

```r
# Fit the joint recurrent + terminal event model (small settings for illustration)
fit <- fit_causal_recur(
  data      = df_fit,
  K         = K,
  id_col    = "id",
  time_col  = "k_fac",
  treat_col = "Ak",
  lag_col   = "lagYk",
  formula_T = Tk ~ Ak + I(lagYk^2) + L.1 + L.2,
  formula_Y = Yk ~ Ak + I(lagYk^2) + L.1 + L.2,
  cores     = 1,
  verbose   = FALSE
)
```

Causal effect estimation via g-computation

Finally, we estimate causal effects corresponding to different treatment initiation strategies using Bayesian
g-computation. In this example, we compare initiation at two different time points. The output summarizes
posterior draws of the causal estimand, including point estimates and uncertainty intervals.

```r
# Bayesian g-computation for two treatment-start strategies
gcomp <- g_computation(
  fit_out = fit,
  s_vec   = c(3, 6),    # start at time 3 vs 6
  B       = 20,
  cores   = 1
)

print(gcomp)
```

4

This quickstart demonstrates the full workflow of BayCauRETM, from data preparation and model fitting to causal effect estimation. Detailed usage and example results are available on GitHub (see the demo PDF).

Acknowledgements

References

Carpenter, Bob, Andrew Gelman, Matthew D. Hoffman, Daniel Lee, Ben Goodrich, Michael Betancourt, Marcus Brubaker, Jiqiang Guo, Peter Li, and Allen Riddell. 2017. "Stan: A Probabilistic Programming Language." *Journal of Statistical Software* 76 (1): 1–32. https://doi.org/10.18637/jss.v076.i01.

Ghosh, Debashis, and D. Y. Lin. 2002. "Marginal Regression Models for Recurrent and Terminal Events." *Statistica Sinica*, 663–88.

Janvin, Matias, Jessica G Young, Pål C Ryalen, and Mats J Stensrud. 2024. "Causal Inference with Recurrent and Competing Events." *Lifetime Data Analysis* 30 (1): 59–118. https://doi.org/10.1007/s10985-023-09594-8.

Oganisian, Arman, Anthony Girard, Jon A Steingrimsson, and Patience Moyo. 2024. "A Bayesian Framework for Causal Analysis of Recurrent Events with Timing Misalignment." *Biometrics* 80 (4): ujae145. https://doi.org/10.1093/biomtc/ujae145.

Schaubel, Douglas E, and Min Zhang. 2010. "Estimating Treatment Effects on the Marginal Recurrent Event Mean in the Presence of a Terminating Event." *Lifetime Data Analysis* 16 (4): 451–77. https://doi.org/10.1007/s10985-009-9149-x.